

# ML Assignment Group: 119

## Problem Statement

**Predict whether the credit card using the customer is going to default or not.**

- Import the data from the [default of credit card clients](https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients) (<https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>) (2 points)
- Consider all columns as independent variables and assign to variable X except the last column and consider the last column as dependent variable and assign to variable y. Remove columns which don't help the problem statement. (1 point)
- Compute some basic statistical details like percentile, mean, standard deviation of dataset (1 point)
- Do Feature Scaling on Independent variables (2 points)
- Split the data into train and test dataset (1 point)
- Use sklearn library to train on train dataset on random forest and predict on test dataset (3 points)
- Compute the accuracy and confusion matrix. (2 points)

## Contributors

J Manoj Balaji (2019AB04228@wilp.bits-pilani.ac.in)

Siddharth Verma (2019AD04097@wilp.bits-pilani.ac.in)

Randeep Singh (2019AD04069@wilp.bits-pilani.ac.in)

## Import necessary dependencies ¶

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
%matplotlib inline
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## Import the dataset

Dropping ID variable since that's a unique variable for each row and doesn't help in prediction


```
In [2]: df = pd.read_excel('default of credit card clients.xls',skiprows=1,usecols="B:Y")
```

```
In [3]: df.head()
```

Out[3]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	...
0	20000	2	2	1	24	2	2	-1	-1	-2	...
1	120000	2	2	2	26	-1	2	0	0	0	...
2	90000	2	2	2	34	0	0	0	0	0	...
3	50000	2	2	1	37	0	0	0	0	0	...
4	50000	1	2	1	57	-1	0	-1	0	0	...

5 rows × 24 columns



## Understanding data


Computing some basic statistical details like percentile, mean, standard deviation of dataset

```
In [4]: df.describe()
```

```
Out[4]:
```

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	3
mean	167484.322667	1.603733	1.853133	1.551867	35.485500	-0.016700	
std	129747.661567	0.489129	0.790349	0.521970	9.217904	1.123802	
min	10000.000000	1.000000	0.000000	0.000000	21.000000	-2.000000	
25%	50000.000000	1.000000	1.000000	1.000000	28.000000	-1.000000	
50%	140000.000000	2.000000	2.000000	2.000000	34.000000	0.000000	
75%	240000.000000	2.000000	2.000000	2.000000	41.000000	0.000000	
max	1000000.000000	2.000000	6.000000	3.000000	79.000000	8.000000	

8 rows × 24 columns



```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   LIMIT_BAL                             30000 non-null  int64
1   SEX                                    30000 non-null  int64
2   EDUCATION                             30000 non-null  int64
3   MARRIAGE                              30000 non-null  int64
4   AGE                                    30000 non-null  int64
5   PAY_0                                 30000 non-null  int64
6   PAY_2                                 30000 non-null  int64
7   PAY_3                                 30000 non-null  int64
8   PAY_4                                 30000 non-null  int64
9   PAY_5                                 30000 non-null  int64
10  PAY_6                                 30000 non-null  int64
11  BILL_AMT1                             30000 non-null  int64
12  BILL_AMT2                             30000 non-null  int64
13  BILL_AMT3                             30000 non-null  int64
14  BILL_AMT4                             30000 non-null  int64
15  BILL_AMT5                             30000 non-null  int64
16  BILL_AMT6                             30000 non-null  int64
17  PAY_AMT1                              30000 non-null  int64
18  PAY_AMT2                              30000 non-null  int64
19  PAY_AMT3                              30000 non-null  int64
20  PAY_AMT4                              30000 non-null  int64
21  PAY_AMT5                              30000 non-null  int64
22  PAY_AMT6                              30000 non-null  int64
23  default payment next month            30000 non-null  int64
dtypes: int64(24)
memory usage: 5.5 MB
```

From the dataset, we can see that SEX, EDUCATION, MARRIAGE, AGE, PAY\_0, PAY\_2, PAY\_3, PAY\_4, PAY\_5, PAY\_6 are categorical variables and BILL\_AMT1, BILL\_AMT2, BILL\_AMT3, BILL\_AMT4, BILL\_AMT5, BILL\_AMT6, PAY\_AMT1, PAY\_AMT2, PAY\_AMT3, PAY\_AMT4, PAY\_AMT5, PAY\_AMT6, LIMIT\_BAL are continuous variables. And as specified in the problem statement, the last column i.e. default payment next month is the target column. We need to scale the values of continuous variables.

From the description mentioned on the dataset repository page [default of credit card clients](https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients). (<https://archive.ics.uci.edu/ml/datasets/default+of+credit+card+clients>) We assign the following list of continuous variables


```
In [6]: continuous_list = ['BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5',  
                           'BILL_AMT6', 'PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6', 'LIMIT_BAL']
```

```
In [7]: categorical_list = np.setdiff1d(df.columns, continuous_list)  
categorical_list
```

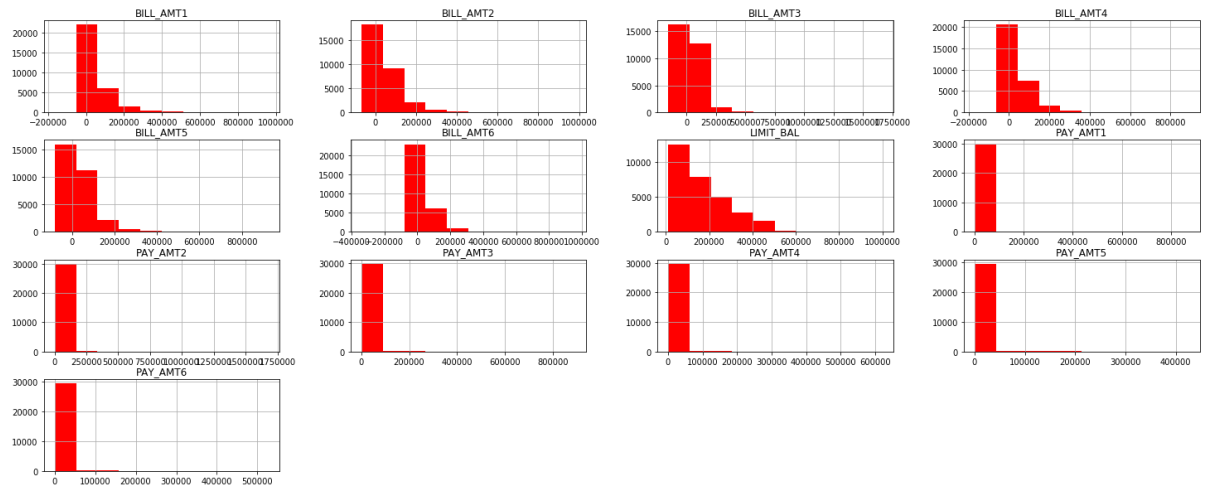
```
Out[7]: array(['AGE', 'EDUCATION', 'MARRIAGE', 'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4',  
              'PAY_5', 'PAY_6', 'SEX', 'default payment next month'],  
             dtype=object)
```

```
In [8]: # understand age to decide whether to consider it as Continuous or Categorical, since the values are within a very short range, we can take it as categorical variable  
sns.distplot(df.AGE, color='tab:orange')  
print('Unique age values:', df.AGE.nunique())
```

Unique age values: 56




```
In [9]: df[continous_list].hist(bins=10, figsize=(25, 10), color='red');
```



```
In [10]: fig, ax = plt.subplots(figsize=(30,15))
sns.heatmap(df.corr(), annot=True, ax=ax, annot_kws={"size": 15})
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1aad6033c18>
```



**Deleting the below attributes because of high correlation amongsts themselves**

All BILL\_AMT variables have high positive correlation amongsts themselves

Pay 3 - 6 have high correlation amongsts themselves


```
In [11]: continous_list
```

```
Out[11]: ['BILL_AMT1',  
          'BILL_AMT2',  
          'BILL_AMT3',  
          'BILL_AMT4',  
          'BILL_AMT5',  
          'BILL_AMT6',  
          'PAY_AMT1',  
          'PAY_AMT2',  
          'PAY_AMT3',  
          'PAY_AMT4',  
          'PAY_AMT5',  
          'PAY_AMT6',  
          'LIMIT_BAL']
```

```
In [12]: removable_continuous = ['BILL_AMT3' , 'BILL_AMT2' , 'BILL_AMT1', 'BILL_AMT4', 'BILL  
_AMT5' , 'BILL_AMT6']  
removable_categorical = ['PAY_3' , 'PAY_4' , 'PAY_5' , 'PAY_6']  
df.drop(removable_continuous, axis=1, inplace=True)  
df.drop(removable_categorical, axis=1, inplace=True)  
continous_list = np.asarray(continous_list)  
categorical_list = np.asarray(categorical_list)  
continous_list = np.setdiff1d(continous_list, removable_continuous)  
categorical_list = np.setdiff1d(categorical_list, removable_categorical)  
print('continous list:',continous_list)  
print('categorical list:',categorical_list)
```

```
continous list: ['LIMIT_BAL' 'PAY_AMT1' 'PAY_AMT2' 'PAY_AMT3' 'PAY_AMT4' 'PAY  
_AMT5'  
'PAY_AMT6']  
categorical list: ['AGE' 'EDUCATION' 'MARRIAGE' 'PAY_0' 'PAY_2' 'SEX'  
'default payment next month']
```

```
In [13]: plt.subplots(figsize=(20,5))
plt.subplot(1,3,1)
ax = sns.countplot(df['SEX'], palette = 'magma')
ax.set_xticklabels(['male', 'female'])
plt.subplot(1,3,2)
ax = sns.countplot(df['MARRIAGE'], palette = 'inferno')
ax.set_xticklabels(['married', 'single', 'others', 'NA'])
plt.subplot(1,3,3)
ax = sns.countplot(df['EDUCATION'].map({1:'graduate school',2:'university',3:
'high school',4:'others',5:'others', 6:'others'}), palette = 'plasma')
temp = ax.set_xticklabels(['graduate school','university','high school','other
s'], rotation=45)
```



## Data Preprocessing

### Feature Scaling on Independent variables

```
In [14]: scaler = StandardScaler()
```

```
In [15]: continous_transform_df = pd.DataFrame(scaler.fit_transform(df[continous_list
]),columns=continous_list)
```

```
In [16]: continous_transform_df.head()
```

Out[16]:

	LIMIT_BAL	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6
0	-1.136720	-0.341942	-0.227086	-0.296801	-0.308063	-0.314136	-0.293382
1	-0.365981	-0.341942	-0.213588	-0.240005	-0.244230	-0.314136	-0.180878
2	-0.597202	-0.250292	-0.191887	-0.240005	-0.244230	-0.248683	-0.012122
3	-0.905498	-0.221191	-0.169361	-0.228645	-0.237846	-0.244166	-0.237130
4	-0.905498	-0.221191	1.335034	0.271165	0.266434	-0.269039	-0.255187

```
In [17]: categorical_df = df[categorical_list]
```

```
In [18]: categorical_df.head()
```

Out[18]:

	AGE	EDUCATION	MARRIAGE	PAY_0	PAY_2	SEX	default payment next month
0	24	2	1	2	2	2	1
1	26	2	2	-1	2	2	1
2	34	2	2	0	0	2	0
3	37	2	1	0	0	2	0
4	57	2	1	-1	0	1	0

```
In [19]: # re-join scaled-continuous and categorical dataset
data_set = categorical_df.merge(continous_transform_df,left_index=True,right_index=True)
```

```
In [20]: data_set.head()
```

Out[20]:

	AGE	EDUCATION	MARRIAGE	PAY_0	PAY_2	SEX	default payment next month	LIMIT_BAL	PAY_AMT1	PAY_AMT2
0	24	2	1	2	2	2	1	-1.136720	-0.341942	-0.221191
1	26	2	2	-1	2	2	1	-0.365981	-0.341942	-0.221191
2	34	2	2	0	0	2	0	-0.597202	-0.250292	-0.136720
3	37	2	1	0	0	2	0	-0.905498	-0.221191	-0.136720
4	57	2	1	-1	0	1	0	-0.905498	-0.221191	1.367202

## Split dataset to train and test datasets

```
In [21]: X=data_set.drop('default payment next month',axis=1)
y=data_set['default payment next month']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [22]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[22]: ((24000, 13), (6000, 13), (24000,), (6000,))

## Use sklearn library to train random forest model on train dataset

```
In [23]: rf = RandomForestClassifier(random_state=42)
```



```
In [24]: rf.fit(X_train,y_train)
```

```
Out[24]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=42, verbose
                                =0,
                                warm_start=False)
```

## Evaluate the model on test dataset

```
In [25]: y_predict=rf.predict(X_test)
```

## Compute the accuracy and confusion matrix and classification report

```
In [26]: accuracy = accuracy_score(y_pred = y_predict,y_true = y_test)
```

```
In [27]: print("Accuracy: {}".format(round(100*accuracy, 4)))
```

Accuracy: 81.4%

```
In [28]: conf_mat = confusion_matrix(y_pred = y_predict,y_true = y_test)
```

```
In [29]: conf_mat
```

```
Out[29]: array([[4425, 262],
                [ 854, 459]], dtype=int64)
```

```
In [30]: print('Classification Report: ')
print(classification_report(y_test, y_predict))
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.94	0.89	4687
1	0.64	0.35	0.45	1313
accuracy			0.81	6000
macro avg	0.74	0.65	0.67	6000
weighted avg	0.79	0.81	0.79	6000

To check the imbalance of the target variable

```
In [31]: Counter(y_test)
```

```
Out[31]: Counter({0: 4687, 1: 1313})
```

## Visualize the Confusion matrix as a heatmap

```
In [32]: fig, ax = plt.subplots(figsize=(7,5))          # Sample figsize in inches
sns.heatmap(conf_mat, annot=True, ax=ax, cmap='Blues', annot_kws={"size": 25},
            fmt="d")
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1aad7fa1b00>
```

